

★ Overview

Agentic Code QA Instruction Guide

Last updated: May 16, 2026

Overview

Agentic Code QA evaluates whether models can answer realistic software-engineering questions by exploring an actual repository at a fixed commit. The project is designed to test **codebase-grounded reasoning**: understanding implementation details, tracing behavior across files, identifying risks or root causes, and explaining conclusions clearly enough for another engineer to trust. The term “QA” in this context is short for “Question & Answer.”

Your work will contribute to **improving AI capabilities for real-world software engineering**. By creating tasks that reflect how engineers actually investigate unfamiliar codebases, you help identify where AI agents still struggle with long-context repository understanding, technical judgment, and reliable developer-facing explanations.

Experts create tasks that **force the model to use the repository** rather than general programming knowledge. For each assigned repo at a specific commit, language, difficulty, and intent, contributors write a concise developer-style prompt (a question), a human-authored golden answer grounded in repository exploration, and a rubric that distinguishes correct, partial, and weak answers.

The rest of this guide contains the concrete working instructions: common mistakes to avoid, LLM usage boundaries, task creation and scope requirements, and reference examples.

Project Expectations

Project Expectations

Writers are expected to produce clear, accurate, and unique tasks while staying active and responsive throughout the project.

Expectation	Details
Quality and Consistency	• Deliver work that is specific to the assigned repository at a specific commit, intent, language, and difficulty. • Keep each task meaningfully distinct from prior submissions; repeated prompt patterns or generic rubrics may be sent back.
Minimum Activity	• Complete at least 5 tasks per week to remain active on the project. • Notify Team Leads in advance if you cannot meet the weekly minimum.
Task Timing	• The client target is for each task to complete the full write, review, and revision cycle within 5 hours. • Aim to write the initial task in 3-4 hours.
Engagement and Communication	• Monitor Slack and task interfaces regularly for reviewer feedback, process updates, and Team Lead messages. • Respond promptly when clarification is needed.



Payments

Payments

Writers are paid per merged task. Reviewers are paid hourly. Promotions may be available for top performers and will be communicated in Slack.

① Base Writer Pay

Base pay depends on how many revision cycles the task needs before it is merged.

- 0-3 revision cycles: \$120 per merged task.
- 4 or more revision cycles: \$0 per merged task. You may still submit the task if you choose, but the payout amount is zero.

② One-Shot Bonuses

One-shot bonuses apply only to tasks that merge with 0 revisions. Bonus amount depends on assigned task difficulty.

- Easy or medium difficulty: +\$30.
- Hard difficulty: +\$50.

③ LLM Tool Reimbursement

Mercor offers a \$20 monthly reimbursement for basic LLM tools, such as Cursor, used to examine repositories. This reimbursement is available after your first task is merged.

- Submit your receipt through the [reimbursement form](#). (Feel free to redact PII)
- Use only tools allowed by the [LLM Usage Policy](#) and project instructions.
- Submit one form per calendar month.

④ Payment Schedule

Payments are calculated on Saturdays (Pacific time), for the previous 7-day Friday -> Friday window. Merged tasks that achieved Merged status *within the previous window*

are included in that week's calculation. Payments typically appear for experts within the following week, per [Mercor payment policies](#).