



Reviewer Guidelines

Reviewer Guidelines

This section outlines responsibilities and guidelines for reviewers. Writers may also read it to understand what reviewers are checking before approval.

New Reviewer Onboarding

Reviewers should be familiar with the full writer workflow, especially Task Claiming & Setup, Repo Exploration & Task Scope, Prompt & Golden Answer Writing, Rubric Writing & Difficulty Calibration, Rubric Design 101, and the Task Examples.

Before reviewing independently, read the examples and ask in the reviewer Slack channel when a task raises a guideline question or ambiguous edge case.

Understanding the Job

Reviewers are guardians of quality. Your primary job as a reviewer is to maintain quality standards while keeping the task pipeline moving. At all times you should operate with professional courtesy, and expect the same from all colleagues. Secondary priorities include:

- Providing education for new writers
- Flagging problems in documentation, workflows, etc.
- Improving your own efficiency (and sharing these improvements)

Reviewer performance is measured by AHT (Average Handling Time) per review, as well as quality (typically measured by overturn rate by Super Reviewer). Work with Team Leads to understand your performance and ensure you meet targets.

Reviewer Guidelines

1. How to Review

Reviewers are responsible for checking every part of the task before approval: prompt, golden answer, rubric, repository evidence, metadata, QC outputs, and difficulty calibration. Use the checklist below to keep reviews consistent.

Review Area	What To Check	Send Back / Escalate If
Assignment &	Repo, commit, language, intent,	Wrong repo/commit/language,

repo setup	difficulty, Dockerfile, and task metadata all match the assignment.	inaccessible repo, broken Dockerfile, or metadata that prevents a fair review.
Prompt quality	Prompt is developer-style, repo-specific, fair, concise, and aligned to the assigned intent/difficulty. It gives starting context without leaking the answer. It reads as human.	Prompt is generic, too broad, a local lookup, answer-leaking, mismatched to intent, non-human sounding (“robotic”), or impossible to verify from the repo.
Golden answer	Answer is correct, complete, grounded in repository evidence, and explains the reasoning chain with relevant files/functions/behaviors. Answer scores 100% on QC.	Unsupported claims, missed causal steps, wrong file/function references, incomplete PR impact, or answer content that the repo does not support.
Rubric quality	Criteria are atomic, measurable, tagged/weighted correctly, include logical dependencies (if any), include fair alternatives, and let the golden answer earn full credit.	Vague criteria, overlapping criteria, placeholder text, direct reversal penalties, style as critical/penalty, or missing important failure modes.
QC & validation	Writer used Prompt QC, Rubric QC, Golden Answer QC, Validate Golden, and Step 5 eval as appropriate, then applied human judgment.	Golden validation fails, QC flags valid unresolved issues, or the task was submitted without the required calibration evidence.
Difficulty & policy	Step 5 score and trajectory fit the assigned band; work follows LLM usage policy and is meaningfully distinct from prior tasks.	Difficulty is out of band, model succeeds with shallow exploration, suspected LLM-authored prompt/golden answer/rubric (non-ASCII characters are a red flag), duplicate task pattern, or policy concern.

2. Reviewer Decision Rules

Use the decision rules below to decide whether to approve, make a small direct edit, send the task back, or escalate it.

Decision	Use When	Reviewer Action
Approve	The task has no substantive issues, or only tiny issues that do not affect task meaning or grading.	Approve. Leave concise feedback if there is a useful note for the writer.
Direct edit + approve	The only issue is low-risk: typo, formatting, obvious label cleanup, or a minor rubric wording fix that does not change task substance.	Make the edit directly, re-run eval, then approve and mention what you changed.
Send back or escalate	Prompt, golden answer, rubric, repo evidence, policy, or difficulty calibration needs substantive writer judgment; or the task is on its 3rd+ revision cycle with unresolved issues.	Send back with specific artifact-level feedback, or escalate to a Team Lead when access/metadata/policy/repeated-cycle issues block a normal review.

3. Priorities

Prioritize review work before writing tasks. Clear reviews from oldest to newest, give priority to writers who may be blocked (new/at claim cap), then review newly submitted tasks. Only after these steps, write new tasks if you are also contributing as a writer. (Do not run your reviewer timer when writing your own tasks.)

4. Communication

Leave feedback that tells the writer exactly what to change and why. Reference the relevant artifact, such as prompt, golden answer, specific rubric criteria, Step 5 eval output, or the task metadata.

Use the reviewer Slack channel for guideline questions, suspected policy violations, repeated writer issues, duplicate task patterns, or project-wide clarifications. Reviewers should proactively answer writer questions when the instruction doc already resolves the issue.

If the documentation is unclear, flag the issue for a Team Lead in Slack. Drive all decisions into documentation, avoiding “lore” stored in chat.

5. Time Expectations

Most reviews should take 30-45 minutes. Spend longer only when the task is hard, repo evidence is subtle, the Step 5 calibration needs careful interpretation, or you are a new reviewer. If a review is blocked by missing access, wrong metadata, or an unusable repo, escalate instead of spending review time reconstructing the task. Rely on Team Leads to flag tasks or handle special cases.

6. Maximum Review Cycles

If a task has already been through 3 or more revision cycles, avoid another broad send-back. Make low-risk direct edits if the task is otherwise approvable, or escalate to a Team Lead when the task still has substantive prompt, answer, rubric, or calibration problems.

7. Exceptions

When a delivery is imminent, sometimes reviewers enter “fix-it mode.” When instructed by a Team Lead, reviewers may prioritize fixing a task (even if handling time increases) in order to meet delivery deadlines.